K2-Think Hackathon API Doc

Overview

The **K2-Think Model API** provides an **OpenAI-compatible endpoint** for querying the MBZUAI-IFM/K2-Think large language model.

This API can be used in the Hackathon environment to interact with the model using standard OpenAI SDKs and tools.

Basic Information

Field	Value
Model Name	MBZUAI-IFM/K2-Think
Base URL	https://llm-api.k2think.ai/v1/
Authentication	Bearer token (API Key distributed to each team)
Rate Limit	30 requests per minute
JSON Schema / Tool Calls	Not supported
Streaming	Supported ("stream": true)

Curl Example

Code block

```
curl -X 'POST' \
'https://llm-api.k2think.ai/v1/chat/completions' \
-H 'accept: application/json' \
-H 'Authorization: Bearer <Your K2 think API Key>' \
-H 'Content-Type: application/json' \
-d '{
"model": "MBZUAI-IFM/K2-Think",
"messages": [
{
"role": "user",
"content": "hello"
}
],
"stream": true
}'
```

Python Example

Code block

```python
# !pip install openai
import argparse
import os
from openai import OpenAI
MODEL_NAME = "MBZUAI-IFM/K2-Think"
DEFAULT_QUESTION = "Solve the 24 game [2,3,5,6]"
def main():
    # Parse command-line arguments
    parser = argparse.ArgumentParser(
        description="Query a model via OpenAI-compatible API")
    parser.add_argument("--question",
                        help="Question to ask the model",
                        default=DEFAULT_QUESTION)
    args = parser.parse_args()
    # Initialize the client with custom endpoint and API key
    client = OpenAI(
        base_url="https://llm-api.k2think.ai/v1",
        api_key="<Your K2 think API Key>",
        timeout=1200
    )
    # Define the conversation, don't include system message
    messages = [
        {"role": "user", "content": args.question}
    ]
    response = client.chat.completions.create(
        model=MODEL_NAME,
        messages=messages,
        stream=False
    )
    print(response.choices[0].message.content)
if __name__ == "__main__":
    main()
```

Disclaimer

- The **K2-Think API** is provided **exclusively for use in this Hackathon**.
- **Do not share your API key** or endpoint configuration outside your team.
- The model may occasionally generate **inaccurate or unsafe outputs**; use discretion when interpreting results.